

W19 Predictive Analysis

What are the difference between fit line? Purposes?

Accurately predict sales price → optimization, analysis

What account for the Change in MSE between test and validation set?

Sampling

```
DataFrame.sample( n=None , frac=None , replace=False)
```

- n= #of items from axis to return
- replace allow sampling from same row more than once
- ignore_index = True : resulting index 0,1,2...n-1

Question Framing

- **relationship** of the variable (being predicted) with other variables

```
#DATA CLEANING
sales=pd.read_csv("NYC_Sales_Samples.csv").replace("-", np.nan).dropna()

sales.iloc[:,2:]=sales.iloc[:,2:].astype("float")
sales[sales.columns[2:]] =sales[sales.columns[2:]].astype("float")

sales=sales.assign(log_GSF=np.log(sales.GROSS_SQUARE_FEET), log_SP=np.log(sales.SALE_PRICE))
#RANDOM SAMPLING
sales=sales.sample(n=sales.shape[0],replace=False,ignore_index=True)
# n= #of items from axis to return
# replace allow sampling from same row more than once
# ignore_index = True : resulting index 0,1,2...n-1
#manipulate existing columns and assign them to table
plt.scatter(sales.log_GSF, sales.log_SP)
plt.xlabel("log GROSS_SQUARE_FEET")
plt.ylabel("log SALE_PRICE")

# fit line

def fit_line(intercept,slope):
    x=np.array([np.min(sales.log_GSF),np.max(sales.log_GSF)])
    y=intercept+slope*x
    plt.plot(x,y, color="red")
fit_line(4,1)
fit_line(5,1)

x=np.linspace(5,13,100)
beta0_g1,beta1_g1=1,2
y=beta0_g1+beta1_g1*x
plt.plot(x,y, label="line1")
beta0_g2,beta1_g2=6,1
y=beta0_g2+beta1_g2*x
plt.plot(x,y, label="line2")
```

1. **Nearest Neighbor** :

2. **Linear regression** : fit a straight line for the dots to be as close as possible

a. Good when outcome variable is continuous

$$y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \epsilon$$

Handwritten annotations:
- y : dependent variable
- $\beta_0, \beta_1, \beta_2$: value to be estimated
- x_1, x_2 : var. that explain variation. (PREDICTORS)
- ϵ : unobserved factor affect y, assume $\epsilon=0$

Variable = predictor + predictor

$$\log_S P = \beta_0 + \beta_1 \log_G SP + \epsilon$$

- Find out the best estimate for β_0 and $\beta_1 \rightarrow$ estimate $y \rightarrow$ predicting

Ordinary Linear Square(OLS)

Find out parameter values($\beta_0, \beta_1, \beta_2$) to **minimize the mean squared errors (MSE)**

$$\operatorname{argmin}_{\beta_0, \beta_1} \frac{1}{N} \sum_{i=1}^N \epsilon_i^2 = \frac{1}{N} \sum_{i=1}^N (y_i - (\beta_0 + \beta_1 x_{1i}))^2$$

$$\hat{y} = \hat{\beta}_0 + \hat{\beta}_1 x_1 + \hat{\beta}_2 x_2$$

Handwritten annotation: estimation of parameters

estimated mean of y

▼ **This can be done using :**

```
sklearn.linear_model.LinearRegression  
model=LinearRegression(fit_intercept=True)
```

- fit_intercept = True : calculate β_0

1. model fitting

`model.fit(X, y)` :

- X is 2D structure. Muti col: `Xsup[:, [0, 3]]`
- y is 1D, `Xsup[:, i]` i is **dependent**
- estimation of the non-intercept β_i : `model.coef_`
 - β_i estimate = change in y when I change 1 unit, other constant
- estimate of intercept β_0 : `model.intercept_`
 - expected value of y when all other variable = 0

```
model.fit(Sales[["log_GSG"]], Sale["log_SP"])
```

2. Prediction

- `model.predict(X)`
 - X = variable which the change in outcome depend on
 - `y = model.intercept_ + model.coef_[0] * X1`

```
prediction=model.predict(sales[["log_GSF"]])
sales=sales.assign(predicted_value=prediction)

plt.plot(sales.log_GSF, sales.predicted_value, label="line1")
```

3. Error estimation (MSE)

`mean_squared_error(y, y_predicted)`


```
prediction=model.predict(sales[["log_GSF"]])
mean_squared_error(sales.log_SP, prediction)
```

Model Selection and Variance & Bias trade off


- Choosing between different models : CROSS VALIDATION (balance between simple and complex)
 - **Underfitting** : too much “Bias”, missing relationship between feature and target
 - simple: high bias, low variance

- **Overfitting:** too much “variance”, model predict noise in training set
 - More complex model: low bias, high variance
 - model sensitive to change in X - lots of noise - hard to predict data never seen before
→ MSE high

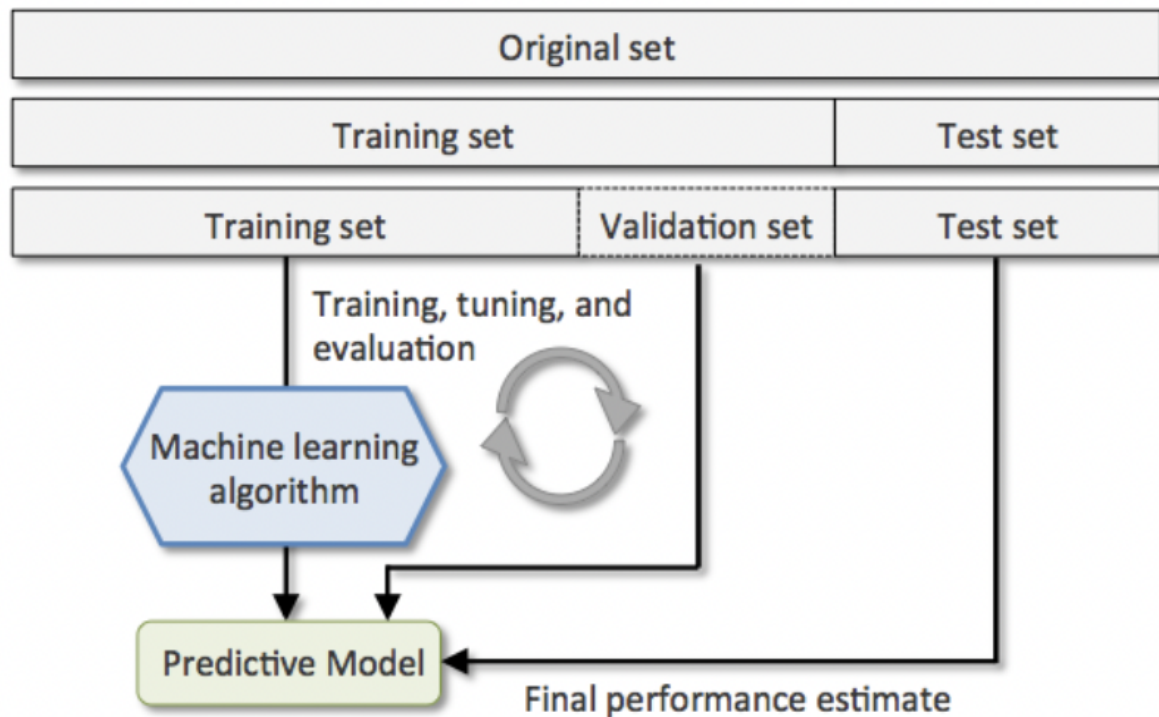
Overfitting: use training data to test



- $y = \beta_0 + \beta_1 x_1 + \epsilon$
- $y = \beta_0 + \beta_1 x_1 + \beta_2 x_1^2 + \epsilon$
- $y = \beta_0 + \beta_1 x_1 + \beta_2 x_1^2 + \beta_3 x_1^3 + \beta_4 x_1^4 + \beta_5 x_1^5 + \beta_6 x_1^6 + \epsilon$



▼ Cross Validation



Training set : loop through candidate model, train with training set (recover **beta** β value)

Validation set : compute accuracy of each model → **choose** the best model

Test set : report **accuracy** of best model

- 1) know best model 2) know accuracy of each model

▼ **OPTIONAL : Combine Train Validation to retrain Best Model:**

```
Train_Valid = pd.concat([Train,Valid])
model.fit(Train_Valid[["log_GSF"]],Train_Valid["log_SP"])
predicted=model.predict(Test[["log_GSF"]])
mean_squared_error(predicted, Test["log_SP"])
```

- Use Test to measure accuracy

▼ **Error Measurement (MSE)**

$$MSE = \frac{\sum (y_i - \hat{y}_i)^2}{N}$$

```
from sklearn.metrics import mean_squared_error
mean_squared_error(y_true, y_predict)
```

Example:

- candidate model
 - Model 1: log(SALES_PRICE) on log(GROSS_SQUARE_FEET)
 - Model 2: log(SALES_PRICE) on log(GROSS_SQUARE_FEET) and total units

```
Train=sales[0:1000:]
Valid=sales[1000:1200:]
Test=sales[1200:]
sales.columns

#training
model1 = LinearRegression()
model1.fit(Train[["log_GSF", "TOTAL_UNITS"]], Train["log_SP"])

model2 = LinearRegression()
model2.fit(Train[["log_GSF"]], Train["log_SP"])

## model selection
model1_predict = model1.predict(Valid[["log_GSF", "TOTAL_UNITS"]])
model2_predict = model2.predict(Valid[["log_GSF"]])
mean_squared_error(Valid.log_SP,model1_predict), mean_squared_error(Valid.log_SP,model2_predict)
```